

Senior Machine Learning Engineer / Research Scientist

Location: Boston, MA (Hybrid)

Department: Artificial Intelligence & Data Science

Employment Type: Full-Time

About the Role

We are seeking a highly skilled Senior Machine Learning Engineer / Research Scientist to join our AI Research & Development team. The ideal candidate will bring deep expertise in machine learning, deep learning, natural language processing, and computer vision, combined with strong software engineering skills. You will lead the design, development, and deployment of cutting-edge ML models, mentor junior researchers, and contribute to innovative projects that bridge advanced research with impactful real-world applications.

Key Responsibilities

- Research & Development: Design, train, and evaluate advanced machine learning models, with a focus on deep reinforcement learning, anomaly detection, and large-scale data processing.
- End-to-End ML Pipelines: Architect and implement scalable, cloud-based ML pipelines for ingesting, processing, and analyzing terabytes of structured and unstructured data.
- Applied AI: Translate complex research findings into production-ready systems for domains such as sensor data analytics, NLP, and computer vision.
- Collaboration & Mentorship: Work closely with cross-functional teams to integrate ML solutions into products; mentor junior engineers and researchers in model development best practices.
- Publications & Knowledge Sharing: Contribute to internal ML libraries and author peer-reviewed papers for leading AI conferences and journals.
- Innovation: Stay up-to-date with the latest ML frameworks, architectures, and algorithms, and propose innovative approaches to solve challenging business and technical problems.

Qualifications

- Education:
 - Ph.D. in Electrical Engineering, Computer Science, or related field (preferred), or equivalent practical experience.
- Strong academic background in AI/ML, reinforced by published research.
- Technical Skills:
 - Programming: Go, Python, C++, Java, R.
 - ML Frameworks: TensorFlow, PyTorch, Keras, scikit-learn.
 - Cloud Platforms: Google Cloud Platform (GCP), AWS, Azure (basic familiarity).
 - Big Data: Apache Spark, Hadoop.
 - Databases: SQL, NoSQL (MongoDB, Cassandra).
 - DevOps: Docker, Kubernetes, Git, Linux.
- Experience:
 - 5+ years designing and deploying ML solutions in production environments.
 - Proven track record in anomaly detection, NLP, computer vision, and reinforcement learning.

- Experience implementing large-scale data processing pipelines.
- Strong publication history in top-tier AI/ML conferences or journals.

Preferred Attributes

- Ability to bridge the gap between cutting-edge research and practical engineering.
- Excellent problem-solving and analytical thinking skills.
- Strong communication and collaboration abilities.
- Passion for innovation and staying ahead in the rapidly evolving AI landscape.