

Maeve Stratton

(617) 555-1212 | maeve.stratton@email.com | linkedin.com/in/maevestratton
| github.com/maevestratton

Professional Summary

Highly motivated and results-oriented Machine Learning Research Fellow at Veridian Dynamics with 5+ years of experience in developing and deploying cutting-edge AI/ML solutions. Proven ability to translate complex research into practical applications, leveraging expertise in deep learning, natural language processing, and computer vision. Seeking a challenging role where I can contribute my skills and experience to drive innovation and solve complex problems. Proficient in Go, Python, C++, TensorFlow, PyTorch, and Google Cloud Platform (GCP).

Education

Massachusetts Institute of Technology (MIT), Cambridge, MA * **Doctor of Philosophy (Ph.D.), Electrical Engineering and Computer Science** | May 2018 * Dissertation: "Novel Architectures for Robust Deep Reinforcement Learning in Uncertain Environments" * GPA: 3.9/4.0 * **Master of Science (M.S.), Electrical Engineering and Computer Science** | May 2016 * **Bachelor of Science (B.S.), Electrical Engineering and Computer Science** | May 2014 * Minor: Mathematics

Technical Skills

Programming Languages: Go, Python, C++, Java, R

Machine Learning Frameworks: TensorFlow, PyTorch, Keras, scikit-learn

Cloud Platforms: Google Cloud Platform (GCP) (Compute Engine, Cloud Storage, BigQuery, Dataflow), AWS (S3, EC2), Azure (basic familiarity)

Databases: SQL, NoSQL (MongoDB, Cassandra)

Big Data Technologies: Apache Spark, Hadoop

Other: Docker, Kubernetes, Git, Linux, Agile methodologies

Professional Experience

Machine Learning Research Fellow | Veridian Dynamics | Boston, MA | June 2020 – Present

- Led the development of a novel deep learning model for anomaly detection in large-scale sensor data, resulting in a 25% improvement in accuracy compared to existing methods.
- Designed and implemented a scalable data pipeline using GCP for processing and analyzing terabytes of data daily.
- Published two peer-reviewed papers on the results of my research at leading AI conferences.
- Mentored junior researchers in best practices for machine learning model development and deployment.
- Contributed to the development of a new internal machine learning library for efficient model training and deployment.

Software Engineer | Innovatech Solutions | Cambridge, MA | June 2018 – June 2020

- Developed and maintained several key features for a large-scale data analytics platform using Python and Java.
- Implemented a new real-time data processing pipeline using Apache Spark, resulting in a 40% reduction in processing time.
- Collaborated with a team of engineers to design and implement a new microservices architecture for improved scalability and maintainability.

AI/ML Projects

Anomaly Detection in Sensor Data (Veridian Dynamics): Developed a deep learning model using a novel architecture based on recurrent neural networks and autoencoders to detect anomalies in large-scale sensor data from industrial equipment. Achieved state-of-the-art performance on a benchmark dataset. (Details available upon request)

Natural Language Processing for Sentiment Analysis (MIT): Built a deep learning model for sentiment analysis of social media text using LSTM networks and attention mechanisms. Improved accuracy by 15% compared to baseline models. (Code available on GitHub)

Computer Vision for Object Detection (Personal Project): Developed a real-time object detection system using YOLOv5 and deployed it on a Raspberry Pi. (Code available on GitHub)

Reinforcement Learning for Robotics (MIT): Developed a reinforcement learning algorithm for controlling a robotic arm to perform complex manipulation tasks in a simulated environment. (Details available upon request)